03. Data Definition using SQL

[PM Jat, DAIICT, Gandhinagar]

Data Definition Language in SQL

What do we mean by data definition?

It is process of defining database schema. In RDBMS, it is done using DDL part of SQL. A data definition program is often referred as DDL script. A DDL script typically contains instructions for creating new database schema.

A data definition task has DDL script as input to DBMS, and on successful run of DDL script on DBMS, following things are done-

- New database instance (empty) is created as per the schema definition given in script.
- Database schema information is saved as data in database catalogue.

SQL provides a set of commands for this purpose and are referred as DDL set of commands of SQL.

Given below is DDL script for XIT database.

```
CREATE TABLE department (
        DID CHAR(2) PRIMARY KEY,
        DNAME VARCHAR(20) NOT NULL
);

CREATE TABLE program (
        PID CHAR(3),
        PNAME VARCHAR(20) NOT NULL,
        INTAKE SMALLINT,
        DID CHAR(2) REFERENCES department(did)
                ON DELETE SET NULL ON UPDATE CASCADE,
        PRIMARY KEY (PID)
);
```

SQL-DDL commands

- This set of commands are used to *define* and *maintain* database schemas
- Following are commonly used commands -
  - CREATE/DROP SCHEMA
  - CREATE/ALTER/DROP TABLE
  - CREATE/DROP SCHEMA
  - CREATE/DROP DOMAIN/TYPE
  - CREATE/DROP VIEWS
- Note that SQL is not case sensitive language.

CREATE/DROP SCHEMA

CREATE SCHEMA command creates a container to hold all elements of a "database instance"; like tables, views, constraints, stored functions, indexes, and so on.

If you use MySQL [Schema=Database] PostgreSQL

For each database, you create a new schema.

Creating new (empty) schema

CREATE SCHEMA XIT;

Drop (existing) schema

DROP SCHEMA XIT;
DROP SCHEMA XIT CASCADE;

Before doing any work on a database, set the database (schema) as "working schema" as following-

set search_path to XIT;

CREATE TABLE statement

`CREATE TABLE` is most comprehensive command for creating relations. This command is used to specify a new relation; command typically requires specifying following parameters

- Name of relation
- Its attributes, domain for attributes, and
- Constraints.

This command does following-

- Creates empty relation in database instance.
- Makes appropriate entries in database catalogue; that is schema definition is put into database catalogue

`CREATE TABLE` – General Form

CREATE TABLE r (A$_1$ dom(A$_1$), Aa$_2$ dom(A$_2$), ..., A$_n$ dom(A$_n$),
(integrity-constraint$_1$),
...,
(integrity-constraint$_k$))

- r is the name of the relation

- each A_i is an attribute name in the schema of relation r

- dom(A_i) is the data type(domain) for attribute A_i

note that in SQL-DDL domains are specified by SQL data types

Example CREATE TABLE

```
CREATE TABLE department (
    DID CHAR(2) PRIMARY KEY,
    DNAME VARCHAR(30) NOT NULL
);

CREATE TABLE program (
    PID CHAR(3),
    PNAME VARCHAR(20) NOT NULL,
    INTAKE SMALLINT,
    DID CHAR(2) REFERENCES department(did),
    PRIMARY KEY (PID)
);

CREATE TABLE student (
    StudID CHAR(3),
    Name VARCHAR(20) NOT NULL,
    ProgID CHAR(3) REFERENCES program(pid),
    "class" smallint,
    cpi decimal(4,2)
);
```

Note: "class" is reserved word is postgresql. Reserved words can be used for attribute names (or for any object for that matter) but need to be put in double quotes, (" "), for example relation student has an attribute named "class".

Default Value [for an attribute]

You can specify default value for an attribute; as shown below

CREATE TABLE EMPLOYEE(

    …

    DNO SMALLINT DEFAULT 4,

    …

)

This value is assigned to the attribute in a new tuple, when we do not specify its value.

Data Types in SQL

DDL uses SQL data types for specifying domain of attributes. Here is broad category of PostgresSQL data types-

- Numeric
- Character strings
- Boolean
- DATE/TIME
- Binary

Numeric Data types

Whole Numbers:

- INTEGER, SMALLINT, BIGINT – though implementation dependent, are typically 2 byte, 4 byte, 8 byte integers

Approximate (Floating point) Numbers:

- REAL, DOUBLE PRECISION – typically 4 and 8 byte floating numbers with precision of 6 and 15 digits precision

- FLOAT(p), lets you specify precision in terms of binary digits; it can be anything between 1 to 53; specifying float with no p is taken as double precision.
- Values are rounded off if have higher precision; Value might overflow and under flow.
- Special values associated with: Infinity, - Infinity, NaN (Not-a-Number)

Exact fractional numbers

- Given inaccuracies with floating point numbers, SQL provides an alternative data type Numeric, where accuracy is important for example in case of Amount, other quantities or so. You would not want Rs 2000 to be shown up as 1999.9999.
- Numeric data types can store numbers with up to 1000 digits of precision and perform calculations exactly.
- However, arithmetic on numeric values is slow compared to the integer types, or to the floating-point types.
- Are declared as DECIMAL(i,j), DEC(i,j) or NUMERIC(i, j), Numeric(i); where i is precision, and j is scale (default scale is 0)
- Example: value 23.2134 has precision 6 and scale 4. If no scale is specified then it is taken as zero

Character Strings

Postgresql provides following three character data types. First two types require specifying the length of string (in terms of number of characters) while third one does not.

Name	Description
character varying(*n*), varchar(*n*)	variable-length with limit
character(*n*), char(*n*)	fixed-length, blank padded
Text	variable unlimited length

First two types can store strings up to n characters in length. An attempt to store a longer string into a column of these types will result in an error, unless excess characters are all spaces, in which case the string will be truncated to the maximum length.

Fixed length characters are typically used for IDs/Codes. Extra spaces are padded towards the end if actual data length is less than specified length.

While character and character varying have maximum length limit as 1GB (take note; this is storage size and not length of string); text data type allows storing strings of unlimited.

As per postgres documentation (as following) -

> There is no performance difference among these three types, apart from increased storage space when using the blank-padded type, and a few extra CPU cycles to check the length when storing into a length-constrained column. While character(n) has performance advantages in some other database systems, there is no such advantage in PostgreSQL; in fact character(n) is usually the slowest of the three because of its additional storage costs.
>
> In most situations text or character varying should be used instead.

Some RDBMS, like oracle have a data type called Character Large Objects (CLOB) having limit of 4GB. [Oracle has limits of 2000 bytes and 4000 bytes on character and character varying respectively]

Boolean

Valid values for Boolean type are

- True: TRUE, 'TRUE'
- False: FALSE, 'FALSE'

Values are not case sensitive

In some implementations (like postgreSQL) you may also find 'y', 'yes', 1 are accepted for true and 'n', 'no', and 0 for false.

Binary data types in SQL

BIT(n): fixed length of n bits

BIT VARYING(n): variable length bits of max n

BYTEA: to store large binary objects like images and video etc.; this is equivalent to Binary Large Objects (BLOB) in some implementations like oracle.

DATE data types

- Stores year, month, and day
- Literals are typically specified as DATE 'YYYY-MM-DD'; for example DATE '1989-05-13'

Note: This is casting syntax as well in ANSI SQL <data-type> data_in_other_type ; in example above string data is gets casted to date type

TIME data types

- Stores hours, minutes, and seconds
- Literals are specified as TIME 'HH:MI:SS', seconds may contain fractional values
- For example- TIME '15:25:18.34'

TIMESTAMP data types

- Stores Y,M,D,H, Min, Second
- Literals are specified as
  TIMESTAMP 'YYYY-MM-DD HH:MI:SS';
  for example-TIMESTAMP '2004-10-19 10:23:54'
- Time and Timestamps can be specified timezone info too, in the form of +13:00 to –12:59 in the format of HH:MI; for example-
  TIMESTAMP '2004-10-19 10:23:54+02'

Specifying Constraints in SQL-DDL

NOT NULL

UNIQUE

PRIMARY KEY

FOREIGN KEY

CHECK

Normally attribute level constraints (involves only one attribute) are specified inline against the attribute definition, while

Constraints involving multiple attributes has to be specified as separate fragment (separated by comma) in CREATE table statement.

Attribute level constraints are in-lined with attribute declaration

Example below:

```
CREATE TABLE PROGRAM (
    PNAME VARCHAR(20) NOT NULL,
    PID SMALLINT PRIMARY KEY,
    DID CHAR(2) NOT NULL REFERENCES
                        DEPARTMENT(DID),
    UNIQUE (PNAME)
);
```

Constraints could be specified as separate fragment (normally done when involves multiple attributes). Example below shows the same.

```
CREATE TABLE PROGRAM (
    PNAME VARCHAR(20) NOT NULL,
    PID SMALLINT,
    DID CHAR(2) NOT NULL,
    UNIQUE (PNAME),
    PRIMARY KEY (PID),
    FOREIGN KEY (DID) REFERENCES
                        DEPARTMENT(DID),
);
```

CHECK constraints:

Two example below should be self-explanatory and conveying the intuition.

```
CREATE TABLE STUDENT (
    …
    CPI REAL CHECK ( cpi >=0 AND cpi <= 10)
    …
);
```

```
CREATE TABLE COURSE_OFFERING (
    …
    semester varchar(6) CHECK ( semester in ('Autumn', 'Winter', 'Summer'))
    …
);
```

Constraints have name in schema

Below is a snap from pgAdmin-III access of "XIT" schema

You may not always be giving names to constraints. PostgreSQL automatically gives names and follows certain conventions while giving names. In constraints shown in above snapshot, you should able to observe following patterns-

PK: <rel-name>_pkey, for example- program_pkey

FK: <rel-name>_<fk-attr-name>_fkey -> <referenced-rel-name>
for example- program_did_fkey->deaprtment

You can give names to constraints explicitly, and should be done for complex constraints, so that you can refer them for dropping or altering purposes

```
CREATE TABLE  works_on (
      essn  CHAR(9),
      ...
      PRIMARY KEY  (essn, pno),
      CONSTRAINT work_proj_check CHECK  ( <predicate> )
)
```

Constraints can be added later too using ALTER TABLE command; example follows.

Create Domain

```
CREATE DOMAIN CPI_TYPE AS REAL CHECK
      (value >= 0 AND value <= 10);
```

This in turn can be used as data type for cpi in student table

```
CREATE TABLE student(
      sid char(9) PRIMARY KEY,
      name varchar(30),
      cpi CPI_TYPE
);
```

DROP TABLE

```
DROP TABLE EMPLOYEE;
```

Drop table not only deletes table rows, but also deletes its definition from the catalog.

Drop table is rejected if table being dropped is referenced by some another relation. To forcefully drop such table using CASCADE option

```
DROP TABLE EMPLOYEE CASCADE;
```

ALTER TABLE

ALTER TABLE EMPLOYEE ADD COLUMN JOB VARCHAR(12);

ALTER TABLE EMPLOYEE DROP COLUMN Address;

If column being deleted is part of some constraints, dropping of column would be rejected; use the cascade option to force the drop, in that case <mark>all constraints using the column will also be dropped</mark>.

ALTER TABLE EMPLOYEE DROP COLUMN Address CASCADE;

Drop default value for an attribute

ALTER TABLE DEPARTMENT ALTER COLUMN MGR_SSN DROP DEFAULT;

Set Default for a column

ALTER TABLE DEPARTMENT ALTER COLUMN MGR_SSN SET DEFAULT '1123456778';

Add a constraint

ALTER TABLE EMPLOYEE
    ADD CONSTRAINT DNO_CHK CHECK (DNO > 0 AND DNO < 21);

Drop a constraint

ALTER TABLE EMPLOYEE DROP CONSTRAINT DNO_CHECK;

General form of alter table

ALTER TABLE <table-name> <action>

Where action can be-

ADD COLUMN

ALTER COLUMN

DROP COLUMN

ADD CONSTRAINT

DROP CONSTRAINT

Referential Actions

- What if we allow deleting tuple with DID='EE' from Department relation?
  Or its code is modified to 'EC'
- Should it be allowed?

Department

DID	DName
CS	Computer Engineering
EE	Electrical Engineering
ME	Mechanical Engineering

Program

ProgID	ProgName	Intake	DID
BCS	BTech(CS)	40	CS
BIT	BTech(IT)	30	CS
BEE	BTech(EE)	40	EE
BME	BTech(ME)	40	ME

- There are some concerns when a tuple referred by some FK value, is being deleted or referred value is being modified
- If this is permitted, the database will be inconsistent; we can specify appropriate course of action in such a situation.
- This is specified by following clause-
  ON UPDATE <action> ON DELETE <action>

And actions could be one of following

- SET NULL – sets any referencing FK to null
- SET DEFAULT - sets any referencing FK to default value, which may be null
- CASCADE – On delete, this deletes any rows with referencing foreign key values. On update this updates, this updates referencing foreign key values to new values.
- NO ACTION- rejects any update that cause referential integrity violation
- RESTRICT- same as NO ACTION with the additional restriction that the integrity check cannot be deferred.

Summary SQL-DDL

CREATE TABLE

ALTER TABLE

CREATE DOMAIN

DROP ...

Add/Drop Constraints

NOT NULL

UNIQUE

PRIMARY KEY

FOREIGN KEY

CHECK

Giving Name to constraints

You can create a table from existing table (or from a result of any SQL query), as following-

```
CREATE TABLE emp_tmp_dno5 AS
      SELECT * FROM employee WHERE dno = 5;
```