

Operations ... SQL Specifics



pm @ daiict



Issues with Null Values

- An attribute having NULL could mean either of following-
 - Value is unknown or not available right now
 - Value is not application for the tuple: a employee not having supervisor will have null in this attribute

- Consider following two SQL statements-

SELECT e.salary*1.1 from employee as e; *NULL*

→ Select * from employee as e where e.salary > 50000; *NULL* *true*

- Interpret e here as tuple variable that ranges over all tuples of employee relations. Try finding result of expressions in blue for tuples where salary is NULL?



How Where clause is used?

- Given a SQL statement

`select * from r`

`where <condition>`

Here r can be any “relational expression”; i.e. `r1 join r2` or so

- You can think of it as following

for tuple $t \in r$

if (cond is true)

add t to resultset





Issues with Null Values

1.1 *e. salary Null

- Arithmetic expressions (+, -, *, /) involving null values result null value for result
- When NULL values appears for attributes used in WHERE clause then boolean expression like this t.a < 10 then interpretation of attribute reference is UNKOWN.
- When we compare a NULL value with another value including NULL, result is UNKNOWN.
- UNKOWN is treated as third truth (in addition to TRUE and FALSE) value in SQL where clause evaluation

e. salary > 50000 and dno = 4 ;
~~UNKNOWN AND TRUE~~ → UNKNOWN
~~FALSE~~



Truth values for UNKNOWN

- NOT
 - NOT UNKNOWN → UNKNOWN
- AND
 - TRUE AND UNKNOWN → UNKNOWN
 - FALSE AND UNKNOWN → FALSE
 - UNKNOWN AND UNKNOWN → UNKNOWN
- OR *ex: salary > 5000 AND OR dno = 4*
 - TRUE OR UNKNOWN → TRUE
 - FALSE OR UNKNOWN → UNKNOWN
 - UNKNOWN OR UNKNOWN → UNKNOWN



Null Values and Comparisons

- While evaluating WHERE clause tuples with UNKNOWN or FALSE truth values are not included in result
- Following query will not include any tuple where either of value ^{is} in NULL irrespective value in other attribute

```
SELECT * FROM EMPLOYEE WHERE  
bdate < DATE '2001-01-01' AND salary > 30000
```

- Following query will not include a tuple only when both are NULL, if one of attribute meets the condition then it will get included in result

```
SELECT * FROM EMPLOYEE WHERE  
bdate < DATE '2001-01-01' OR salary > 30000
```



Null Values and Comparisons

– IS NULL

- Following will not give desired result. Why? -

```
SELECT * FROM employee  
WHERE superssn = NULL;
```

- This is so because Null = Null is also UNKNOWN. For checking an attribute for having NULL value, SQL provides IS NULL (and IS NOT NULL)
- We write as following for such situations –

```
SELECT * FROM employee  
WHERE superssn IS NULL ;
```

IS NOT NULL



Sub-queries in SQL



Subquery in SQL

- A Query that is part of another query is *subquery*. A subquery may also have subquery, and so forth upto any level

Select ✓ (*Select* ✓)

- A subquery in SQL is written as a *query expression* enclosed in parentheses, and is in following form-

" (**SELECT** . . . **FROM** ...) "

as a part of some existing query

- Result of sub-query is again a relation;



Subquery in FROM clause

- Here is an example

```
select ename, dno, salary from employee natural join  
  (select eno from works_on natural join project  
   where pno=1) as r;
```

```
select r1.*, m.salary as manager_salary from  
  (select e.eno as enum, e.ename as emp_name,  
         e.salary as emp_salary,  
         d.mgr_eno as manager  
   FROM employee e join department d  
   on e.dno=d.dno) as r2  
  join employee m on (manager=eno)  
 where emp_salary > m.salary;
```



Subquery in FROM clause

```
select eno, ename, salary from employee
natural join
(select eno from employee
except
select mgr_eno from department) as r;
```



Subquery in WHERE clause

```
select * from employee where eno not in  
(select eno from works_on);
```

- Sub query in FROM clause is never a problem. However, when subquery in where clause, there needs to be few cautions to be noted!
- Recall that where clause is evaluated for every tuple of operand relation?
 - Does it mean; subquery here needs to be executed N times?
 - Note subquery here may also be called as inner query!



Execution of Subquery

- **SUB-Query may not execute for every tuple of outer query**

- Consider another query-

```
SELECT * FROM student WHERE  
progid IN (SELECT pid FROM program  
           WHERE did = 'EE' );
```

- Typically, after execution of inner query, outer query may be translated to:

```
SELECT * FROM student WHERE  
progid IN (BEC, BEE) ;
```

- However this optimization may not be possible when you have “*correlated sub-query*”



Correlated Sub-Queries

- When inner query makes a reference to tuple of outer query then it is correlated sub-query. Consider following query -

- List employees, whose salary is more than department average:

```
SELECT eno, ename FROM employee as e
WHERE salary > (SELECT AVG(salary) FROM
employee WHERE dno = e.dno)
```



Execution of Correlated Sub-Queries

- Consider same query

```
SELECT eno, ename FROM employee as e
WHERE salary > (SELECT AVG(salary)
FROM employee WHERE dno = e.dno)
```

- Logically, it is as following: For each tuple of outer query, execute inner query.
- Note that it can not be executed once for all tuples of outer query, as the case be with un-related inner query, and we have to execute **SUB-Query for every tuple of outer query**
- This is identified problem with correlated sub-queries.



Correlated Sub-Queries could be expensive to execute – therefore should be avoided

- Correlated queries are expensive to execute, and can be avoided; for example the previous example
- **SELECT eno, ename FROM employee as e
WHERE salary > (SELECT AVG(salary)
FROM employee WHERE dno = e.dno)**
- can be re-written as-
**SELECT eno, ename, salary FROM employee as e
NATURAL JOIN (SELECT dno, AVG(salary) as
avg_sal FROM employee GROUP BY dno) as av
WHERE salary > av.avg_sal;**



more Correlated Sub-queries

- List down employees having salary greater than their immediate supervisors.

```
select * from employee as e1 where e1.salary >
(select salary from employee as e2 where e2.eno =
e1.super_eno) ;
```

- Select employees having dependents older than 18 years:

```
SELECT * FROM employee AS e WHERE eno IN (SELECT
essn FROM dependent AS d WHERE essn = e.eno AND
age(d.bdate) > interval '18 years') ;
```
- Attempt re-writting them without correlated query.



Compare a values with a bag of values (SQL)

- For example consider following two queries [Find out employee who have salary greater some or all employees of dno = 4]

SELECT enon, ename FROM employee WHERE salary

> **SOME** (SELECT salary FROM employee WHERE dno = 4);

SELECT eno, ename FROM employee WHERE salary

> **ALL** (SELECT salary FROM employee WHERE dno = 4);



Compare a values with a bag of values (SQL)

- Note the equivalences:

SELECT enon, ename FROM employee WHERE salary

> **SOME** (SELECT salary FROM employee WHERE dno = 4); and

SELECT enon, ename FROM employee WHERE salary

> (SELECT **min**(salary) FROM employee WHERE dno = 4);

SELECT enon, ename FROM employee WHERE salary

> **ALL** (SELECT salary FROM employee WHERE dno = 4); and

SELECT enon, ename FROM employee WHERE salary

> (SELECT **max**(salary) FROM employee WHERE dno = 4);



Compare a values with a bag of values (SQL)

- Comparative operators could be, one of following-
>SOME, >=SOME, <=SOME, <SOME, =SOME, <>SOME
>ALL, >=ALL, <=ALL, <ALL, =ALL, <>ALL
- Note: it can be easily proved that
=SOME is identical to **IN**, and
<>SOME is not identical to **NOT IN**
= ALL is not identical to **IN**
<> ALL (mean = NONE) and is same as **NOT IN**, and
Earlier versions of SQL used **ANY** for **SOME**; today both keywords are used as synonymous.



Sub-queries in Update statements

- **UPDATE** employee
 SET salary = salary * 1.1
 WHERE ssn IN (...);
- **DELETE** employee
 WHERE ssn = (...);



EXISTS and NOT EXISTS in SQL

- Checks for emptiness of a relation and returns true or false.
- **EXISTS (r)** can be interpreted as “is there some tuple exists in relation r”
- **EXISTS (r)** returning true says that argument relation r is not empty
- Similarly, **NOT EXISTS (r)** returning true says that argument relation r empty



Example EXISTS

- List employees who have dependents older than 18 years
- `SELECT * FROM employee AS e WHERE EXISTS
(SELECT * FROM dependent AS d WHERE d.eno
= e.eno AND age(d.bdate) > interval '18
years');`



SQL- EXISTS and IN

- While they might appear to be serving similar purposes, semantically are different.
- Both appear as part of predicate in WHERE clause of SELECT
- IN:
 - Syntax: **x IN (r)**
 - Meaning: checks existence of tuple x in relation r, if found returns true, other wise false. Normally x is a scalar value and r is a single column relation.
- EXISTS:
 - Syntax: **EXISTS (r)**
 - Meaning: checks if r is a non empty relation. Returns true if the relation has at least one tuple, otherwise false.
- In both above cases **r** is a *relational expression* resulting a relation.



Bags and Relational Operations



Relational Operations and multiset (or bag)

- By Definition, relations are set; but implementations may permit duplicate tuples and such relations are called *bags*
- Normally stored relations (base) relations should still be sets, because most relations have Primary Key
- However SQL SELECT results are often bags, possibly because duplicate removal is expensive.
- To get *set* you use DISTINCT keyword



SQL and Multiset (or Bag)

- SET operations, that are UNION, INTERSECT, and EXCEPT in SQL yield their result as SET, that means duplicates are removed
- SQL however provides options by which you can have bag results by adding ALL keyword to operation name, i.e. UNION ALL, EXCEPT ALL or so.
- Let us see an example-



UNION/INTERSECT/EXCEPT ALL in SQL

- Compare result of following queries:

```
SELECT superssn FROM employee; --Q1
```

```
SELECT mgrssn FROM department; --Q2
```

```
SELECT superssn FROM employee  
UNION
```

```
SELECT mgrssn FROM department; --Q3
```

```
SELECT superssn FROM employee  
UNION ALL
```

```
SELECT mgrssn FROM department; --Q4
```



R	superssn
	102
	101
	101
	102
	101
	(Null)
	108
	108

S	mgrssn
	101
	102
	108

R UNION S	superssn
	101
	102
	108
	(Null)

R UNION ALL S	superssn
	102
	101
	101
	102
	101
	(Null)
	108
	108
	101
	102
	108



R	superssn
	102
	101
	101
	102
	101
	(Null)
	108
	108

S	mgrssn
	101
	102
	108

R EXCEPT S	superssn
	(Null)

R EXCEPT ALL S	superssn
	101
	101
	102
	108
	(Null)



UNION/INTERSECT/EXCEPT ALL in SQL

- UNION ALL
 - count of an element e in result is sum of count in R and S
- INTERSECT ALL
 - $\min(\text{count-}r, \text{count-}s)$ of an element in R and S , is taken as result
- EXCEPT ALL:
 - Every occurrence of an element e in S decreases its count in R by one.



Functions and Operators in SQL



Functions and Operators

- SQL provides various functions and operators that can be used to create a new attribute in resultant relations
- There are typically, type conversion, arithmetic operators, mathematical, and string manipulation operators and functions. For example: `substring`, `upper`, `lower`, `sqrt`, `ln`, etc.
- Details for PostgreSQL functions can be seen at:
http://intranet.daiict.ac.in/~pm_jat/postgres/html/functions.html.



Examples

```
SELECT ssn,  
       fname || ' ' || minit || '. ' || lname AS name,  
       current_date - bdate AS age FROM employee;
```

```
SELECT essn, hours*50 AS amount FROM works_on;
```

```
SELECT upper(fname) AS name, ln(salary) AS x FROM  
employee;
```

```
SELECT * FROM employee  
       WHERE upper(fname) = 'FRANKLIN';
```

```
SELECT essn FROM dependent WHERE age(d.bdate) > interval  
'18 years');
```



BETWEEN and LIKE in SQL

- **BETWEEN, LIKE** are used in predicate:
 - **SELECT ... WHERE A BETWEEN 10 TO 20;**
 - **SELECT ... WHERE A1 LIKE '%IX%' OR A2 LIKE 'ABC%' OR A3 LIKE '%XYZ';**
 - **SELECT ... WHERE A1 LIKE '_X_%';**
- Also: **NOT BETWEEN** and **NOT LIKE**.



Regular Expression Matching in PostgreSQL

- PostgreSQL also allows regular expression matching in string match using **IS SIMILAR TO <reg-ex>**



ORDER BY

- ORDER BY CLAUSE is used for ordering the resultant tuples of a SQL query.
- Following statements returns row-set from employee table, and rows are sorted based on salary. To order in descending order, we add DESC keyword after attribute name.

```
select * from employee order by salary;
```

```
select * from employee order by salary desc;
```

- Following statement returns row-set from employee table, and rows are sorted in ascending order of dno, and within dno all rows are sorted on salary in descending order-

```
select * from employee order by dno, salary desc;
```



LIMIT and OFFSET

- Examples below should be self explanatory
- Gives top three earners

```
select * from employee order by salary desc limit 3
```

- Gives next two earners after top 3

```
select * from employee order by salary desc  
offset 3 limit 2
```