

Systems
Software/Programming
Device Driver
<https://lwn.net/Kernel/LDD3/>

Linux Architecture

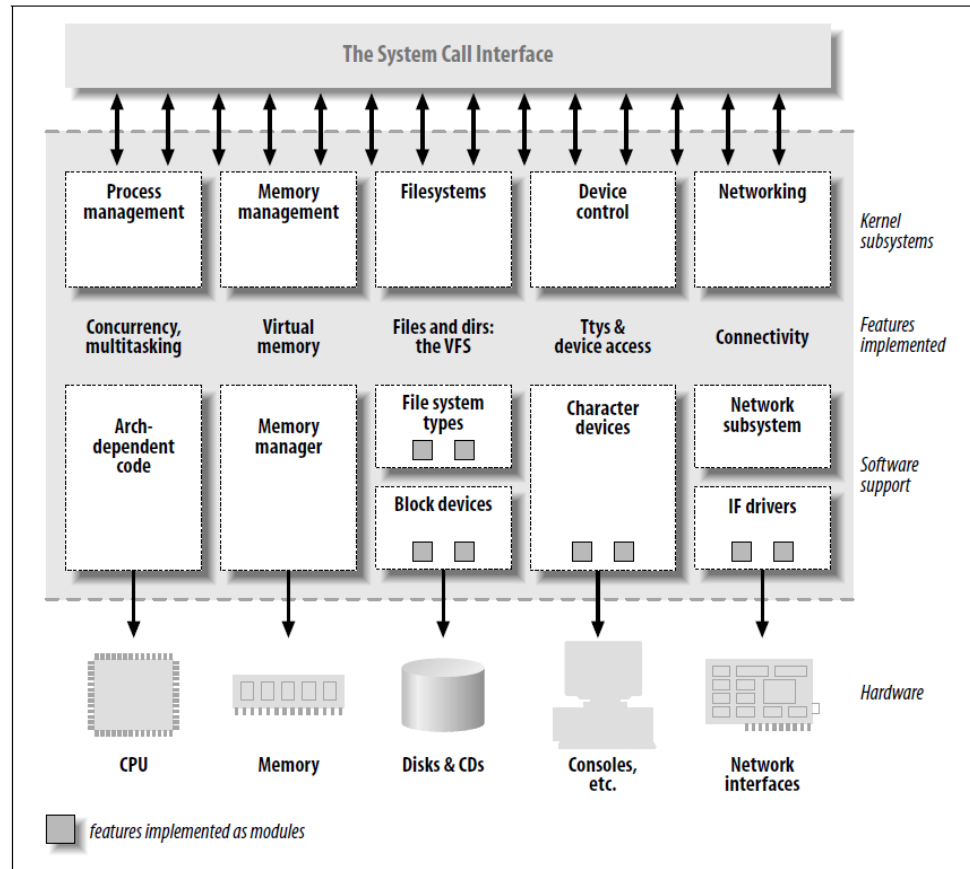


Figure 1-1. A split view of the kernel

- Linux Kernel has several modules:
 - Process Management (studied)
 - Memory Management – create and manage of virtual address space
 - Filesystems – Manages files and directories stored in physical disk by dividing it into logical components.
 - Device Control/Driver – Software designed to communicate and operate specific hardware so that kernel use the hardware without having to know the details.
 - Networking – Protocol implementation

Device Driver Types

- Character Device
 - Read/write character at a time and implements open, read, write, close system calls
 - Character devices are synchronous and only one process can use it at a time
 - Examples: console, keyboard, mouse etc
- Block Device
 - Read/write block of data at a time
 - Writing is done asynchronously so multiple process can write at the same time
 - Examples: CDROM, HDD, USB etc
- Network Device
 - Sends or receives data in form of packets.
 - Example: Hardware device: Ethernet card, Software device: Loopback adapter

Device Driver Module Information

- License:
 - Specified using `MODULE_LICENSE("<license>")` macro
 - `<license>` : "GPL", "GPL v2", "Dual BSD/GPL", "Dual MIT/GPL", "Proprietary" etc
- Author:
 - Specified using `MODULE_AUTHOR("Author")` macro
- Module Description:
 - Specified using `MODULE_DESCRIPTION("My Driver")` macro
- Module Version:
 - Format [`<epoch>:`]`<version>`[`-<extra-version>`]
 - Specified using `MODULE_VERSION("2:1.0")` macro

Constructor and Exit functions

- Init function: `module_init(* func())`
macro

- Called when driver module is loaded
e.g. using `insmod`

```
static int __init mydriver_init(void) /*  
Constructor */  
{  
    return 0;  
}  
module_init(mydriver_init);
```

- Exit function: `module_exit(* func())`
macro

- Called when driver module is
unloaded e.g. using `rmmod`

```
void __exit mydriver_exit(void)  
{  
  
}  
module_exit(mydriver_exit);
```

printk()

- `printk()` kernel level function similar to `printf()` which is user level
- Except, `printk()` allow assigning loglevel/priority of messages
 - `KERN_EMERG`: emergency
 - `KERN_ALERT`: requiring immediate attention
 - `KERN_CRIT`: critical condition related to hardware or software
 - `KERN_ERR`: error reporting
 - `KERN_WARNING`: warning reporting
 - `KERN_NOTICE`: take a note
 - `KERN_INFO`: informational messages
 - `KERN_DEBUG`: debugging information

Simple Driver

- Source [DeviceDriver\mydriver1.c](#) and makefile [DeviceDriver\mydriver1 makefile](#)
- Compile: `$sudo make`
- List module: `$lsmod | grep mydriver1`
- Load module: `$sudo insmod mydriver1.ko`
- Unload module: `$ sudo rmmod mydriver1`
- To see messages from `printk()`: `$dmesg`

Passing Parameters to Driver Code

DeviceDriver\mydriver2.c

DeviceDriver\mydriver2 makefile

- `module_param(name, type, perm)`: Initialize primitive type variable
 - creates the sub-directory under `/sys/module`
 - Type: `bool`, `invbool`, `charp` (char ptr), `int`, `long`, `short`, `uint`, `ulong`, `ushort`
 - Perm: `S_I(W/R)(USR/GRP/OTH)`, can be combined with bit OR (`|`)
- `module_param_array(name, type, num, perm)`: For arrays
- `module_param_cb()`: register callback function which is called when change in value is detected in parameter.
- Load module: `$ $ sudo insmod mydriver2.ko value=13 name="DAIICT" arr_value=10,20,30,40`
- Files for each parameter with its values is created in `/sys/module/mydriver2/parameters/` → “cat name” or “cat cb_value”
- To change the value:
 - First login as root user by `$su`
 - Change value : `$sudo bash -c 'echo new_value > /sys/module/mydriver2/parameters/param_name'`
 - e.g. `sudo bash -c 'echo 20 > /sys/module/mydriver2/parameters/cb_value'`
 - Check messages in system log using `dmesg`

All devices have Major and Minor Numbers

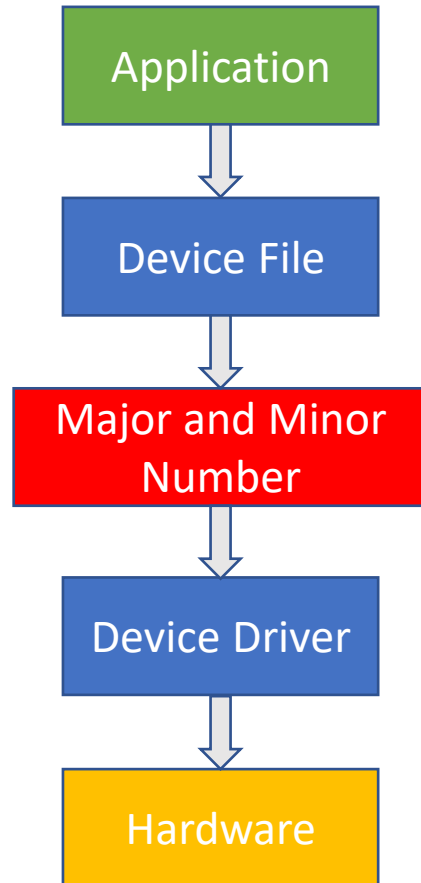
- Major number represents which driver will be used to manage a particular device. It allows multiple devices to have the same major number if they are managed by the same driver.
- Minor number represents a particular device that is managed by the driver identified using Major number.
- In the example below there are 4 tty devices represented as minor numbers 0, 10, 11, 12 but they are all managed by same driver number 4. Similarly 2 disk sda devices with minor numbers 0 and 1 are managed by driver major number 8.

\$ ls -ltr /dev

```
crw--w---- 1 root  tty   4,  0 May 15 12:46 tty0
crw--w---- 1 root  tty   4, 10 May 15 12:46 tty10
crw--w---- 1 root  tty   4, 12 May 15 12:46 tty12
crw--w---- 1 root  tty   4, 11 May 15 12:46 tty11
brw-rw---- 1 root disk 8,  0 May 15 12:46 /dev/sda
brw-rw---- 1 root disk 8,  1 May 15 12:46 /dev/sda1
```

C stands for char device

B stands for block device



Character Device Driver → Create Major and Minor Numbers DeviceDriver\mydriver3.c DeviceDriver\mydriver3 makefile


- `dev_t` struct from `linux/types.h` used to represent `<major>:<minor>`
- Create device's major and minor nos using macro `MKDEV(int major, int minor)` that will return `dev_t` structure.

```
dev_t dev = MKDEV(201, 0)
```

- Statically Register the Device → you provide value for major number, **can have conflict even if different drivers are used.**

```
int register_chrdev_region(dev_t first, unsigned int count, char *name);
```

- `first` → starting device number; `dev_t dev = MKDEV(201, 0)`
- `count` → number of devices requested i.e. number of minor numbers
- `name` → device name representing number range `[first, first+count]`



Character Device Driver → Create Major and Minor Numbers DeviceDriver\mydriver3.c DeviceDriver\mydriver3 makefile

- Dynamically Register the device → Major number is allocated to you by kernel

```
int alloc_chrdev_region(dev_t *dev, unsigned int firstminor,  
unsigned int count, char *name);
```

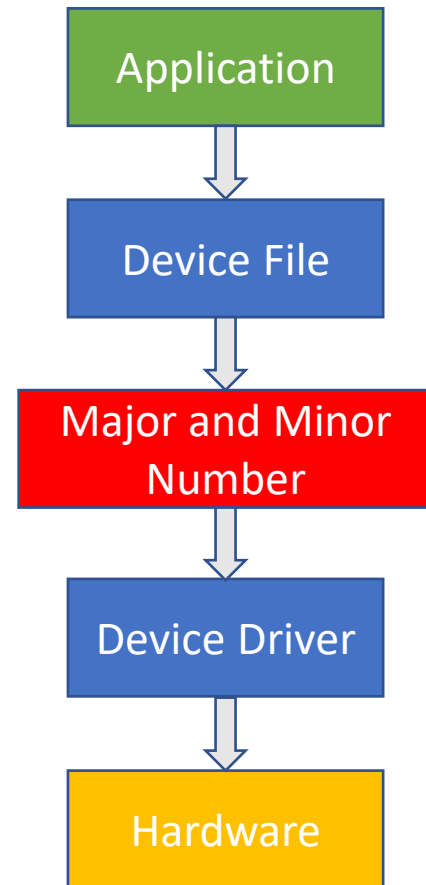
- dev → major number is returned by kernel which is free to use
- Firstminor → you assign starting minor number

- Unregister the device

```
void unregister_chrdev_region(dev_t first, unsigned int count);
```


After driver is loaded using insmod, **check major number for your device in “cat /proc/devices”**

You will not find physical device in /dev directory yet since we have not created the actual device which will be done next.




Character Device Driver → Create Device File

- All devices are treated as files and are stored in /dev directory in Unix-based OS.
- Create the device using shell command mknod:
 - `mknod -m <permissions> <name> <device type> <major> <minor>`
 - name – full path e.g. /dev/mydev
 - device type – c for char device, b for block device
 - Example: `$sudo mknod -m 666 /dev/mydev c 201 0`
 - Check device creation using `→ ls -l /dev/ | grep <devname>`
- **Better approach is to create the device as part of the driver code itself.**



Character Device Driver → Create Device File

- Automatically (Programmatically) [DeviceDriver\mydriver4.c](#)
[DeviceDriver\mydriver4 makefile](#)
 - header file linux/device.h and linux/kdev_t.h
 - Create structure class for our driver under /sys/class:
 - `struct class * class_create (struct module *owner, const char *name);`
 - `void class_destroy (struct class * cls); // destroy once done`
 - Create device:
 - `struct device *device_create (struct class * cls, struct device *parent, dev_t dev, const char *fmt, ...);`
 - `void device_destroy (struct class * class, dev_t devt); // destroy when done`
 - Check device creation using → `ls -l /dev/ | grep <devname>`



cdev Structure linux/cdev.h

- cdev structure is used by linux kernel to represent char device. **Inode of this device will have pointer to cdev structure**

```
struct cdev {  
    struct kobject kobj;  
    struct module *owner;  
    const struct file_operations *ops;  
    struct list_head list;  
    dev_t dev;  
    unsigned int count;  
};
```

- owner = THIS_MODULE defined in linux/module.h
- ops = list of function pointers which implements the operations for this device
- Kobj = information about driver kernel object
- Dev = your device

inode structure has pointer to cdev structure.
In the virtual box installation at /usr/src/linux-headers-5.8.0-48-generic/include/linux/fs.h

```
struct inode {  
    umode_t  
    i_mode;  
    unsigned short  
    i_opflags;  
    kuid_t          i_uid;  
    kgid_t          i_gid;  
    unsigned int    i_flags;  
    .  
    .  
    .  
    union {  
        struct pipe_inode_info  
        *i_pipe;  
        struct block_device    *i_bdev;  
        struct cdev            *i_cdev;  
        char                   *i_link;  
        unsigned               i_dir_seq;  
    };  
};
```

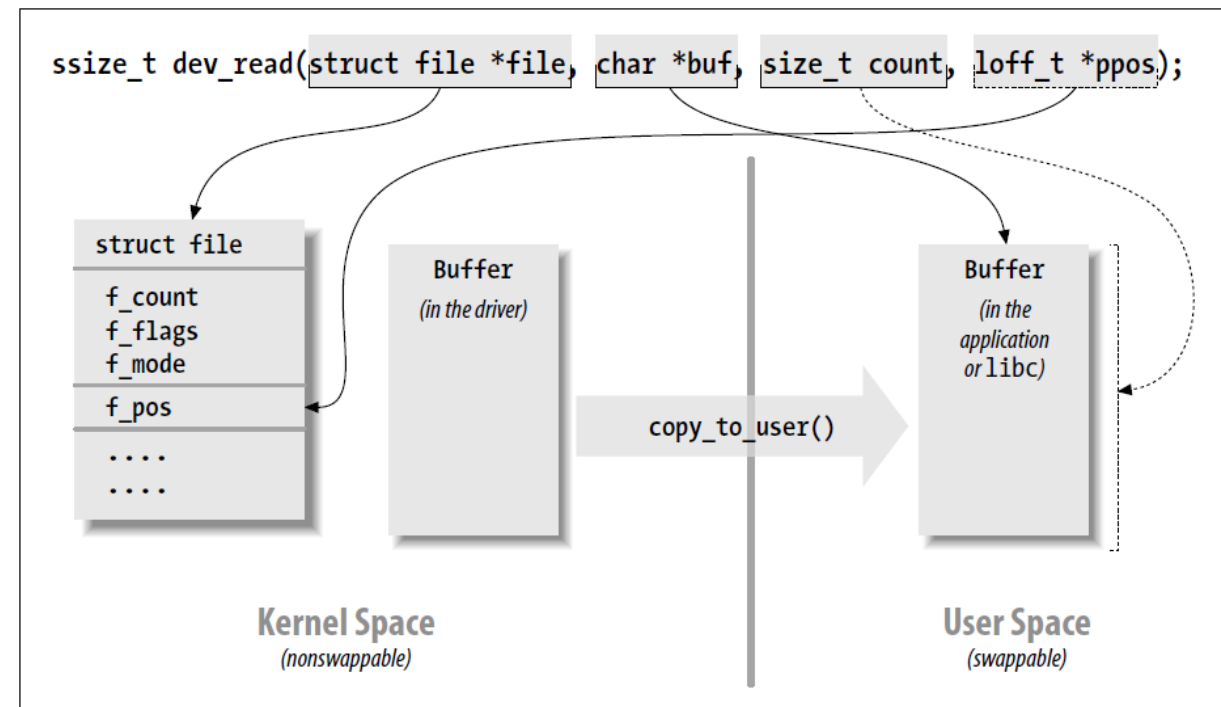
File Operation Structure defined in linux/fs.h

```
struct file_operations mydev_fops = {  
    .owner = THIS_MODULE,  
    .read = mydev_read,  
    .write = mydev_write,  
    .ioctl = mydev_ioctl,  
    .open = mydev_open,  
    .release = mydev_release,  
};
```

There are many other file_operations structure members but have listed important ones only.

File Operation Structure defined in linux/fs.h

- `.owner : struct module *owner` → owner of the structure (THIS_MODULE)
- `mydev_read: ssize_t (*read) (struct file *file, char __user * buf, size_t count, loff_t *ppos);` → function pointer which implements reading of data from device.
 - `char __user *` is pointer to user mode buffer to which data need to be read
 - `size_t` is number of bytes to be read
 - `loff_t` represents current file position for reading and writing.



File Operation Structure defined in linux/fs.h

- `mydev_write: ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *)`; → function pointer which implements writing of the data to device. `loff_t` represents current file position for reading and writing.
- `mydev_ioctl : int (*ioctl) (struct inode *, struct file *, unsigned int, unsigned long)`; → `ioctl` function implementation such as formatting etc
- `mydev_open : int (*open) (struct inode *, struct file *)`; → open the device and allocate file structure
- `mydev_close: int (*release) (struct inode *, struct file *)`; → close device and release file structure

Useful memory function inside driver code

- `void *kmalloc(size_t size, gfp_t flags);` defined in `linux/slab.h`
 - `GFP_USER` – Allocate memory on behalf of user. May sleep.
 - `GFP_KERNEL` – Allocate normal kernel ram. May sleep.
 - `GFP_ATOMIC` – Allocation will not sleep. May use emergency pools. For example, use this inside interrupt handlers.
 - `GFP_HIGHUSER` – Allocate pages from high memory.
 - `GFP_NOIO` – Do not do any I/O at all while trying to get memory.
 - `GFP_NOFS` – Do not make any fs calls while trying to get memory.
 - `GFP_NOWAIT` – Allocation will not sleep.
 - `__GFP_THISNODE` – Allocate node-local memory only.
- `void kfree(const void *objp)` → free allocated memory
- `unsigned long copy_from_user(void *to, const void __user *from, unsigned long n);` → copy data from user space buffer to kernel space buffer
- `unsigned long copy_to_user(const void __user *to, const void *from, unsigned long n);` → copy data from kernel space buffer to user space buffer

Cdev and File Operations

[DeviceDriver\mydriver5.c](#)

[DeviceDriver\mydriver5 makefile](#)

- Create file_operation structure with function pointers for function that you implement
static struct file_operations fops =

```
{  
.owner      = THIS_MODULE,  
.read       = mydev_read,  
.write      = mydev_write,  
.open       = mydev_open,  
.release    = mydev_release,  
};
```

- Initialize cdev structure with specific file_operations

```
void cdev_init(struct cdev *cdev, struct file_operations *fops);
```

- Tell kernel which device uses this cdev structure

```
int cdev_add(struct cdev *cdev, dev_t dev, unsigned int count);
```

Test the driver Code using Application

- Your device will by default have permission to the user who created the device which is generally root because we use sudo

```
$ ls -ltr /dev/mydevice
```

```
crw----- 1 root root 243, 0 Mar 28 15:41 mydevice
```

- Enable the read-write permission for other users too (otherwise only root user can access the device)

```
$ ls -ltr /dev/mydevice
```

```
crw-rw-rw- 1 root root 243, 0 Mar 28 15:41 mydevice
```

Now you can test device with your application

[DeviceDriver\mydriver5 test.c](#)