

# Video Streaming and CDNs: context

- video traffic: major consumer of Internet bandwidth
  - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
  - ~1B YouTube users, ~75M Netflix users
- challenge: **scale** - how to reach ~1B users?
  - single mega-video server won't work (why?)
- challenge: heterogeneity
  - different users have different capabilities (e.g., wired versus **mobile**; bandwidth rich versus bandwidth poor)
- **solution**: distributed, application-level infrastructure




Application Layer 2-2

# Multimedia: video

- video: sequence of images displayed at constant rate
  - e.g., 24 images/sec
- digital image: array of pixels
  - each pixel represented by bits
- coding: use redundancy *within* and *between* images to decrease # bits used to encode image
  - spatial (within image)
  - temporal (from one image to next)

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (*purple*) and number of repeated values ( $N$ )



frame  $i$

*temporal coding example:* instead of sending complete frame at  $i+1$ , send only differences from frame  $i$




frame  $i+1$

# Multimedia: video

- **CBR: (constant bit rate):**  
video encoding rate fixed
- **VBR: (variable bit rate):**  
video encoding rate changes  
as amount of spatial,  
temporal coding changes
- **examples:**
  - MPEG I (CD-ROM) 1.5 Mbps
  - MPEG2 (DVD) 3-6 Mbps
  - MPEG4 (often used in Internet, < 1 Mbps)

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (*purple*) and number of repeated values ( $N$ )



frame  $i$

*temporal coding example:* instead of sending complete frame at  $i+1$ , send only differences from frame  $i$



frame  $i+1$

■ Video → seq. of images shown at fast rate  $\approx 30/s$   
 audio / Image  
 seq of images







Image. / X . d P G  
 pixel → color. (R, G, B)  
 jpeg → image compression algo  
 Image → 3 color images  
 RGB → 8 bit →  $N \times M$  pixels. →  $K$  bits brightness.  
 (0 → 255)



# ■ Compression of Image → Video



Images: Human beings are not able to diff. between fine details  
 → remove some fine details without us noticing it  
 → JPEG → Spatial Fourier transform and throw out high freq components.  
 .BMP → .JPG



+ Compress individual video frames [I]

+ code diff. between seq. of frames.

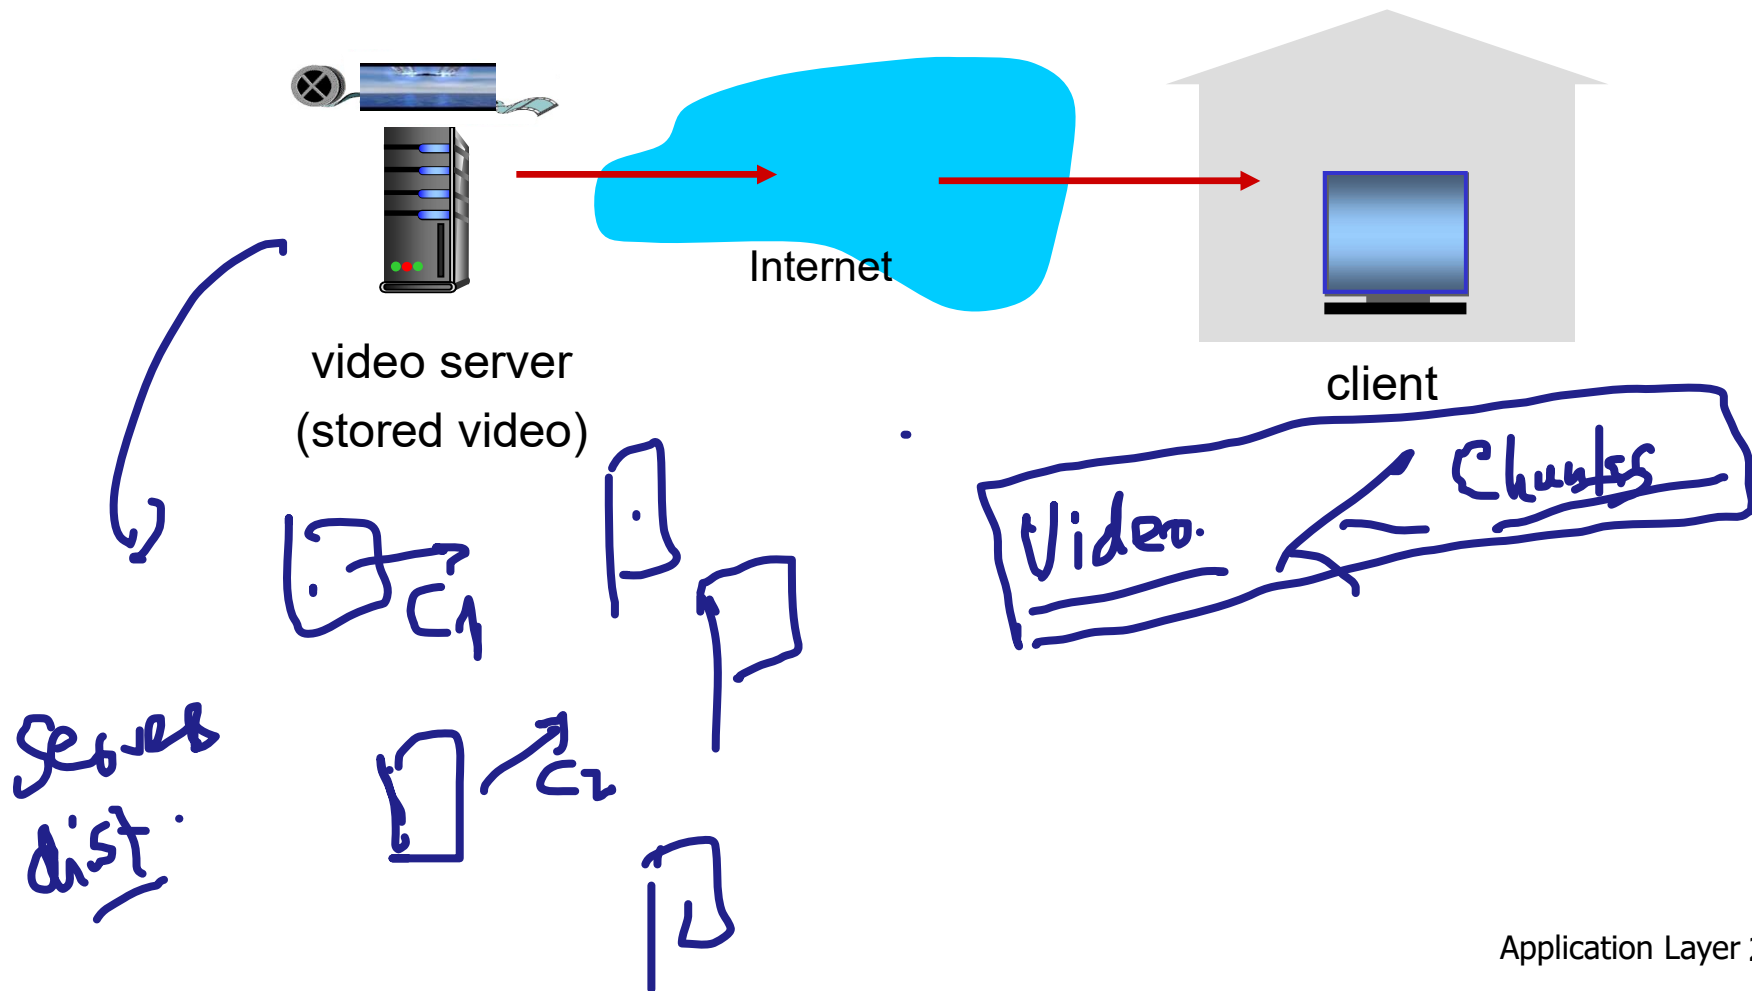
Temporal Correlation

→ very high compression.

coding: MPEG, H.264, H.265

# Streaming stored video:

simple scenario:





manifest file

→ Video Name, metadata

→  $\frac{C1.4K}{C1.HD}$  ———→ { IP addr / port no. (URL) } set

→ C2 ———→ { - - - - - }



# Streaming multimedia: DASH

HDTV -  
4K -  
mobile -  
lower Rq.

- **DASH**: Dynamic, Adaptive Streaming over HTTP

- server:

- divides video file into multiple chunks
- each chunk stored, encoded at different rates
- manifest file provides URLs for different chunks

4K -  
HD -  
mobile -  
lower Rq.

- client:

- periodically measures server-to-client bandwidth
- consulting manifest, requests one chunk at a time
  - chooses maximum coding rate sustainable given current bandwidth → best Rq.
  - can choose different coding rates at different points in time (depending on available bandwidth at time)

best URL


# Streaming multimedia: DASH

---

- *DASH: Dynamic, Adaptive Streaming over HTTP*
- “intelligence” at client: client determines
  - *when* to request chunk (so that *buffer starvation*, or overflow does not occur)
  - *what encoding rate* to request (higher quality when more bandwidth available)
  - *where* to request chunk (can request from URL server that is “close” to client or has high available bandwidth)

# Content distribution networks

- **challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?
- **option 1:** single, large “mega-server”
  - single point of failure
  - point of network congestion
  - long path to distant clients
  - multiple copies of video sent over outgoing link




....quite simply: this solution **doesn't scale**

# Content distribution networks

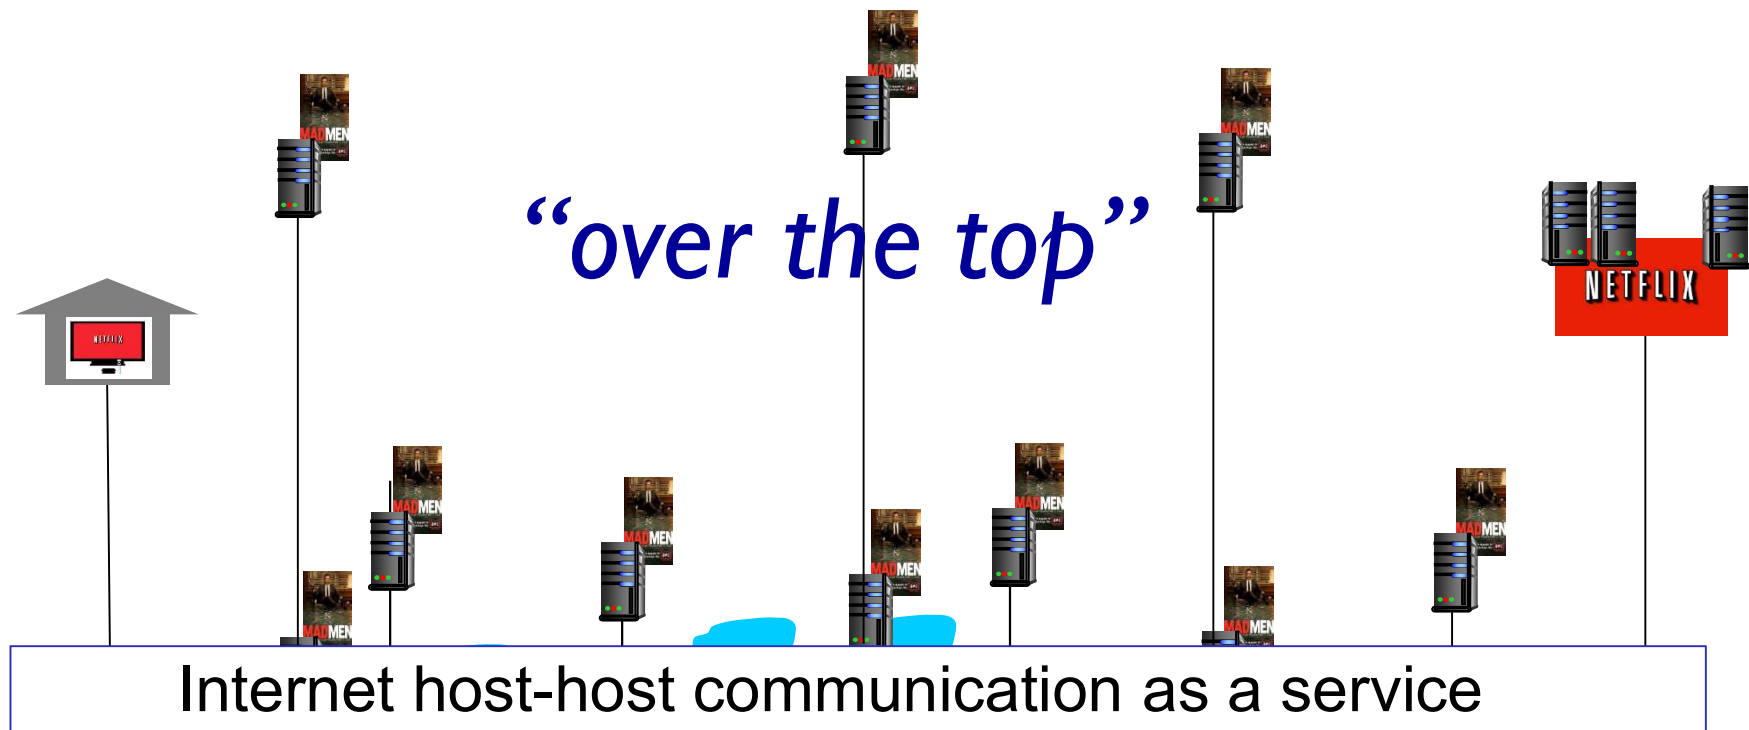
- *challenge*: how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- *option 2*: store/serve multiple copies of videos at multiple geographically distributed sites (*CDN*)
  - *enter deep*: push CDN servers deep into many access networks
    - close to users
    - used by Akamai, 1700 locations
  - *bring home*: smaller number (10's) of larger clusters in POPs near (but not within) access networks
    - used by Limelight

# Content Distribution Networks (CDNs)

- CDN: stores copies of content at CDN nodes
  - e.g. Netflix stores copies of MadMen
- subscriber requests content from CDN
  - directed to nearby copy, retrieves content
  - may choose different copy if network path congested



# Content Distribution Networks (CDNs)



**OTT challenges:** coping with a congested Internet


- from which CDN node to retrieve content?
- viewer behavior in presence of congestion?
- what content to place in which CDN node?

*more .. in chapter 7*

# CDN content access: a closer look

Bob (client) requests video <http://netcinema.com/6Y7B23V>

- video stored in CDN at <http://KingCDN.com/NetC6y&B23V>





# Case study: Netflix

